

notebook

April 13, 2024

```
[1]: import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import seaborn as sns
import numpy as np
from sklearn.inspection import DecisionBoundaryDisplay
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
```

```
[2]: data = pd.read_csv("framingham.csv")
data
```

```
[2]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	\
0	1	39	4.0	0	0.0	0.0	
1	0	46	2.0	0	0.0	0.0	
2	1	48	1.0	1	20.0	0.0	
3	0	61	3.0	1	30.0	0.0	
4	0	46	3.0	1	23.0	0.0	
...	
4235	0	48	2.0	1	20.0	NaN	
4236	0	44	1.0	1	15.0	0.0	
4237	0	52	2.0	0	0.0	0.0	
4238	1	40	3.0	0	0.0	0.0	
4239	0	39	3.0	1	30.0	0.0	

	prevalentStroke	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	\
0	0	0	0	195.0	106.0	70.0	26.97	
1	0	0	0	250.0	121.0	81.0	28.73	
2	0	0	0	245.0	127.5	80.0	25.34	
3	0	1	0	225.0	150.0	95.0	28.58	
4	0	0	0	285.0	130.0	84.0	23.10	
...	
4235	0	0	0	248.0	131.0	72.0	22.00	
4236	0	0	0	210.0	126.5	87.0	19.16	

4237	0	0	0	269.0	133.5	83.0	21.47
4238	0	1	0	185.0	141.0	98.0	25.60
4239	0	0	0	196.0	133.0	86.0	20.91

	heartRate	glucose	TenYearCHD
0	80.0	77.0	0
1	95.0	76.0	0
2	75.0	70.0	0
3	65.0	103.0	1
4	85.0	85.0	0
...
4235	84.0	86.0	0
4236	86.0	NaN	0
4237	80.0	107.0	0
4238	67.0	72.0	0
4239	85.0	80.0	0

[4240 rows x 16 columns]

```
[3]: # check for any null values
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4240 entries, 0 to 4239
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                   4240 non-null   int64
1   age                    4240 non-null   int64
2   education              4135 non-null   float64
3   currentSmoker          4240 non-null   int64
4   cigsPerDay              4211 non-null   float64
5   BPMeds                 4187 non-null   float64
6   prevalentStroke         4240 non-null   int64
7   prevalentHyp            4240 non-null   int64
8   diabetes               4240 non-null   int64
9   totChol                4190 non-null   float64
10  sysBP                  4240 non-null   float64
11  diaBP                  4240 non-null   float64
12  BMI                    4221 non-null   float64
13  heartRate              4239 non-null   float64
14  glucose                 3852 non-null   float64
15  TenYearCHD             4240 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 530.1 KB
```

```
[4]: data.describe()
```

```
[4]:
```

	male	age	education	currentSmoker	cigsPerDay \
count	4240.000000	4240.000000	4135.000000	4240.000000	4211.000000
mean	0.429245	49.580189	1.979444	0.494104	9.005937
std	0.495027	8.572942	1.019791	0.500024	11.922462
min	0.000000	32.000000	1.000000	0.000000	0.000000
25%	0.000000	42.000000	1.000000	0.000000	0.000000
50%	0.000000	49.000000	2.000000	0.000000	0.000000
75%	1.000000	56.000000	3.000000	1.000000	20.000000
max	1.000000	70.000000	4.000000	1.000000	70.000000

	BPMeds	prevalentStroke	prevalentHyp	diabetes	totChol \
count	4187.000000	4240.000000	4240.000000	4240.000000	4190.000000
mean	0.029615	0.005896	0.310613	0.025708	236.699523
std	0.169544	0.076569	0.462799	0.158280	44.591284
min	0.000000	0.000000	0.000000	0.000000	107.000000
25%	0.000000	0.000000	0.000000	0.000000	206.000000
50%	0.000000	0.000000	0.000000	0.000000	234.000000
75%	0.000000	0.000000	1.000000	0.000000	263.000000
max	1.000000	1.000000	1.000000	1.000000	696.000000

	sysBP	diaBP	BMI	heartRate	glucose \
count	4240.000000	4240.000000	4221.000000	4239.000000	3852.000000
mean	132.354599	82.897759	25.800801	75.878981	81.963655
std	22.033300	11.910394	4.079840	12.025348	23.954335
min	83.500000	48.000000	15.540000	44.000000	40.000000
25%	117.000000	75.000000	23.070000	68.000000	71.000000
50%	128.000000	82.000000	25.400000	75.000000	78.000000
75%	144.000000	90.000000	28.040000	83.000000	87.000000
max	295.000000	142.500000	56.800000	143.000000	394.000000

	TenYearCHD
count	4240.000000
mean	0.151887
std	0.358953
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
[5]: # Let's fill all the null values by their mean value
data["education"] = data["education"].fillna(data["education"].mean())
data["cigsPerDay"] = data["cigsPerDay"].fillna(data["cigsPerDay"].mean())
data["BPMeds"] = data["BPMeds"].fillna(data["BPMeds"].mean())
data["totChol"] = data["totChol"].fillna(data["totChol"].mean())
data["BMI"] = data["BMI"].fillna(data["BMI"].mean())
data["heartRate"] = data["heartRate"].fillna(data["heartRate"].mean())
```

```
data["glucose"] = data["glucose"].fillna(data["glucose"].mean())
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4240 entries, 0 to 4239
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4240 non-null   int64
1   age                   4240 non-null   int64
2   education             4240 non-null   float64
3   currentSmoker         4240 non-null   int64
4   cigsPerDay            4240 non-null   float64
5   BPMeds                4240 non-null   float64
6   prevalentStroke       4240 non-null   int64
7   prevalentHyp          4240 non-null   int64
8   diabetes              4240 non-null   int64
9   totChol               4240 non-null   float64
10  sysBP                 4240 non-null   float64
11  diaBP                 4240 non-null   float64
12  BMI                   4240 non-null   float64
13  heartRate             4240 non-null   float64
14  glucose               4240 non-null   float64
15  TenYearCHD            4240 non-null   int64
dtypes: float64(9), int64(7)
memory usage: 530.1 KB
```

```
[6]: # By using scaler , scale the data so that it can be easily analyzed
scaler = StandardScaler()
model = scaler.fit(data)
scaled_data = model.transform(data)
data
```

```
[6]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
0         1   39         4.0              0          0.0  0.000000
1         0   46         2.0              0          0.0  0.000000
2         1   48         1.0              1         20.0  0.000000
3         0   61         3.0              1         30.0  0.000000
4         0   46         3.0              1         23.0  0.000000
...     ...   ...         ...              ...         ...
4235      0   48         2.0              1         20.0  0.029615
4236      0   44         1.0              1         15.0  0.000000
4237      0   52         2.0              0          0.0  0.000000
4238      1   40         3.0              0          0.0  0.000000
4239      0   39         3.0              1         30.0  0.000000


      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
0                   0              0         0     195.0  106.0   70.0  26.97
```

1	0	0	0	250.0	121.0	81.0	28.73
2	0	0	0	245.0	127.5	80.0	25.34
3	0	1	0	225.0	150.0	95.0	28.58
4	0	0	0	285.0	130.0	84.0	23.10
...
4235	0	0	0	248.0	131.0	72.0	22.00
4236	0	0	0	210.0	126.5	87.0	19.16
4237	0	0	0	269.0	133.5	83.0	21.47
4238	0	1	0	185.0	141.0	98.0	25.60
4239	0	0	0	196.0	133.0	86.0	20.91

	heartRate	glucose	TenYearCHD
0	80.0	77.000000	0
1	95.0	76.000000	0
2	75.0	70.000000	0
3	65.0	103.000000	1
4	85.0	85.000000	0
...
4235	84.0	86.000000	0
4236	86.0	81.963655	0
4237	80.0	107.000000	0
4238	67.0	72.000000	0
4239	85.0	80.000000	0

[4240 rows x 16 columns]

```
[7]: # let's check if education affects heart disease
plt.scatter(data["education"],data["TenYearCHD"])
plt.xlabel("education")
plt.ylabel("heart disease")
plt.show()
```



```
[8]: # as seen in the graph, education doesn't show any direct relation to having
      ↪ heart disease. So we can delete that column.
      data.drop(["education"] , axis=1)
```

```
[8]:
```

	male	age	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	\
0	1	39	0	0.0	0.000000	0	
1	0	46	0	0.0	0.000000	0	
2	1	48	1	20.0	0.000000	0	
3	0	61	1	30.0	0.000000	0	
4	0	46	1	23.0	0.000000	0	
...	
4235	0	48	1	20.0	0.029615	0	
4236	0	44	1	15.0	0.000000	0	
4237	0	52	0	0.0	0.000000	0	
4238	1	40	0	0.0	0.000000	0	
4239	0	39	1	30.0	0.000000	0	

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	\
0	0	0	195.0	106.0	70.0	26.97	80.0	
1	0	0	250.0	121.0	81.0	28.73	95.0	
2	0	0	245.0	127.5	80.0	25.34	75.0	

3	1	0	225.0	150.0	95.0	28.58	65.0
4	0	0	285.0	130.0	84.0	23.10	85.0
...
4235	0	0	248.0	131.0	72.0	22.00	84.0
4236	0	0	210.0	126.5	87.0	19.16	86.0
4237	0	0	269.0	133.5	83.0	21.47	80.0
4238	1	0	185.0	141.0	98.0	25.60	67.0
4239	0	0	196.0	133.0	86.0	20.91	85.0

	glucose	TenYearCHD
0	77.000000	0
1	76.000000	0
2	70.000000	0
3	103.000000	1
4	85.000000	0
...
4235	86.000000	0
4236	81.963655	0
4237	107.000000	0
4238	72.000000	0
4239	80.000000	0

[4240 rows x 15 columns]

```
[9]: X = np.
      ↪asarray(data[["male","age","currentSmoker","cigsPerDay","BPMeds","prevalentStroke","prevalen
y = np.asarray(data['TenYearCHD'])
# split the data in two parts as train data and test data in 80% and 20% ratio
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size = 0.3, random_state = 4)

# LogisticRegression
model = LogisticRegression(max_iter=3000)
model.fit(X_train, y_train)

# Prediction
y_pred = model.predict(X_test)

# check for the accuracy
print(accuracy_score(y_pred,y_test))
print(classification_report(y_pred,y_test))
```

0.8553459119496856


	precision	recall	f1-score	support
0	0.99	0.86	0.92	1252
1	0.06	0.60	0.12	20

accuracy			0.86	1272
macro avg	0.53	0.73	0.52	1272
weighted avg	0.98	0.86	0.91	1272

```
[10]: # confusion matrix
cm = confusion_matrix(y_test, y_pred)
conf_matrix = pd.DataFrame(data = cm,
                           columns = ['Predicted:0', 'Predicted:1'],
                           index = ['Actual:0', 'Actual:1'])

plt.figure(figsize = (8, 5))
sns.heatmap(conf_matrix, annot = True, fmt = 'd', cmap = "Greens")

plt.show()
print('The details for confusion matrix is =')
print(classification_report(y_test, y_pred))
```



The details for confusion matrix is =

	precision	recall	f1-score	support
0	0.86	0.99	0.92	1084
1	0.60	0.06	0.12	188

accuracy			0.86	1272
macro avg	0.73	0.53	0.52	1272
weighted avg	0.82	0.86	0.80	1272

The above model which uses a logistic regression model has around 85-86% accuracy. So this model is pretty accurate to make predictions.

Let's try and implement SVM model and check for accuracy.

```
[11]: #Build the model
svm = SVC(kernel="rbf")
# Trained the model
svm.fit(X_train, y_train)
predict_y = svm.predict(X_test)
print(classification_report(predict_y,y_test))
```

	precision	recall	f1-score	support
0	1.00	0.85	0.92	1270
1	0.01	0.50	0.01	2

accuracy			0.85	1272
macro avg	0.50	0.68	0.47	1272
weighted avg	1.00	0.85	0.92	1272

This SVM model also has 85% accuracy same as logistic regression. So both models can be used for making accurate predictions.

Both the models have high accuracy but the f1-score in predicting 1 is very low meaning that both the models are not accurate in predicting the patients who have cardiac disease. However in predicting the patients who do not have cardiac disease, both models are almost 100% accurate.

Let's implement and check accuracy by using Random Forest Algorithm.

```
[18]: rfc = RandomForestClassifier(n_estimators = 500)

rfc.fit(X_train, y_train)

# performing predictions on the test dataset
y_pred = rfc.predict(X_test)

print(accuracy_score(y_pred,y_test))
print(classification_report(y_pred,y_test))
```

```
0.8482704402515723
```

	precision	recall	f1-score	support
0	0.99	0.85	0.92	1255

	1	0.03	0.35	0.06	17
accuracy				0.85	1272
macro avg		0.51	0.60	0.49	1272
weighted avg		0.98	0.85	0.91	1272

Random Forest Classifier also shows around the same results meaning that there is not any significant difference between the accuracy of these three models. And maybe using Logistic Regression for this problem might be a better choice to get more accurate results.